CHATBOT PLATFORM
USING
LMM & RAG

KRITIKA, ANSHDEEP SINGH , GURJOT SINGH

Student UID: 23BCA10554, 23BCA10392, 23BCA10061

DEPARTMENT OF COMPUTING

CHANDIGARH UNIVERSITY

A Thesis/Dissertation

Submitted in Partial Fulfilment of

the Requirements for the Degree of

Bachelors of Computer Applications

in the field of

COMPUTING

SUBMITTED TO:- Mr. Ravinder Singh

DATE OF SUBMISSION: April 15, 2026

ABSTRACT

KRITIKA, GURJOT, ANSHDEEP, for the COMPUTING Degree, presented on April 15, 2026, at

CHANDIGARH UNIVERSITY.

Title: - CHATBOT PLATFORM USING LMM & RAG

Professor: **Mr. Ravinder Singh**

Artificial intelligence has grown rapidly in recent years and has significantly influenced human-computer interaction. Traditional chatbots, which relied on rule-based approaches or limited machine learning techniques, often lacked the ability to understand complex queries, maintain contextual awareness, and provide accurate, real-time responses. With the emergence of advanced neural architectures, especially Large Language Models (LLMs) and Large Multimodal Models (LMMs), chatbot systems have evolved into more sophisticated and human-like conversational agents. However, despite these advancements, significant challenges remain, including hallucination of information, reliance on static training data, lack of transparency, and limited access to up-to-date knowledge sources.

This research paper presents the design and development of a chatbot platform that integrates Large Multimodal Models (LMMs) with Retrieval-Augmented Generation (RAG) to address the limitations of existing systems. The proposed system combines the generative capabilities of LMMs with an external knowledge retrieval mechanism, enabling the chatbot to access dynamic, domain-specific, and up-to-date information. By incorporating RAG, the chatbot retrieves relevant information from a vector database or knowledge repository and uses it as contextual input for generating accurate and meaningful responses. This hybrid approach significantly improves the factual correctness and relevance of responses while reducing the risk of hallucination.

The methodology of the study involves the development of a modular chatbot architecture consisting of a user interface, backend processing unit, retrieval system, and LMM-based response generator. The system processes user queries by converting them into embeddings, retrieving relevant documents using similarity search techniques, and generating context-aware responses through the LMM. Additionally, the platform supports multimodal inputs such as text, images, and documents, enhancing user interaction and making the system more versatile across different application domains.

The proposed chatbot platform is evaluated based on several performance metrics, including response accuracy, latency, contextual relevance, scalability, and user satisfaction. Experimental results indicate that the integration of RAG with LMM significantly enhances the chatbot's ability to provide precise and contextually appropriate answers compared to traditional LLM-based systems. Furthermore, the system demonstrates improved efficiency in handling large datasets and maintaining low response times, making it suitable for real-time applications.

This study also addresses key challenges such as data privacy, system scalability, and efficient knowledge integration. By leveraging vector databases and optimized retrieval techniques, the chatbot ensures secure and efficient data handling while maintaining high performance. The inclusion of multimodal capabilities further extends the applicability of the system to domains such as education, healthcare, customer support, and enterprise solutions.

In conclusion, the integration of Large Multimodal Models with Retrieval-Augmented Generation represents a significant advancement in chatbot technology. The proposed system provides a robust and scalable solution for building intelligent, accurate, and context-aware conversational agents. This research contributes to the growing field of conversational AI by offering a comprehensive framework for developing next-generation chatbot platforms that overcome the limitations of existing approaches and meet the evolving demands of users.

ACKNOWLEDGEMENT

I would like to express my sincere gratitude to all those who supported me throughout the course of this research. I would like to thank **Mr. Ravinder Singh** for his invaluable assistance and insights leading to the writing of this paper at every stage. His insights and expertise have been instrumental in shaping this research.

I would also like to acknowledge **CHANDIGARH UNIVERSITY** for providing the necessary resources and a conducive research environment.

My sincere thanks also goes to the members of my graduate committee for their patience and understanding during the one month of effort that went into the production of this paper.

Lastly, I am grateful to my family and friends for their constant encouragement and patience throughout this journey.

TABLE OF CONTENTS

1. INTRODUCTION
2. METHODOLOGY
3. PROBLEM STATEMENT
4. OBJECTIVES OF THE STUDY
5. JUSTIFICATIONS OF THE STUDY
6. LITERATURE REVIEW
7. SCOPE OF STUDY
8. REFERENCES

INTRODUCTION

Artificial Intelligence has brought major changes in how humans interact with computer systems, particularly through the development of intelligent chatbot systems. Chatbots are software applications designed to simulate human conversation using natural language processing (NLP) techniques. Over time, chatbot technology has evolved from rule-based systems to more advanced AI-driven systems capable of understanding context, intent, and semantics.

The emergence of Large Language Models (LLMs), and more recently Large Multimodal Models (LMMs), has revolutionized chatbot development. LMMs extend traditional language models by incorporating multiple data modalities such as text, images, audio, and video. These models are capable of understanding complex queries, generating human-like responses, and supporting diverse applications such as virtual assistants, customer service, education, and healthcare.

Despite their impressive capabilities, LMMs suffer from certain limitations. One of the major issues is hallucination, where the model generates incorrect or misleading information. Additionally, these models rely on pre-trained knowledge and may lack access to real-time or domain-specific data. This creates challenges in applications that require high accuracy and upto-date information.

To overcome these limitations, Retrieval-Augmented Generation (RAG) has emerged as a powerful technique. RAG enhances the capabilities of LMMs by integrating external knowledge retrieval systems. Instead of relying solely on internal model knowledge, RAG retrieves relevant information from external databases or documents and uses it as context for generating responses. This approach significantly improves accuracy, reduces hallucination, and ensures that responses are grounded in factual data.

The integration of LMM and RAG creates a robust chatbot platform capable of delivering intelligent, context-aware, and reliable responses. Such systems are particularly useful in knowledge-intensive domains where accuracy and relevance are critical.

This research paper aims to explore the design, development, and implementation of a chatbot platform using LMM and RAG. It discusses the methodology, challenges, objectives, and potential applications of such systems while providing insights into current research trends and future directions.

METHODOLOGY

The development of a chatbot platform using LMM and RAG involves a structured and systematic approach. The methodology consists of multiple stages, each contributing to the overall functionality and performance of the system.

Architecture of Chatbot Platform using LMM and RAG :

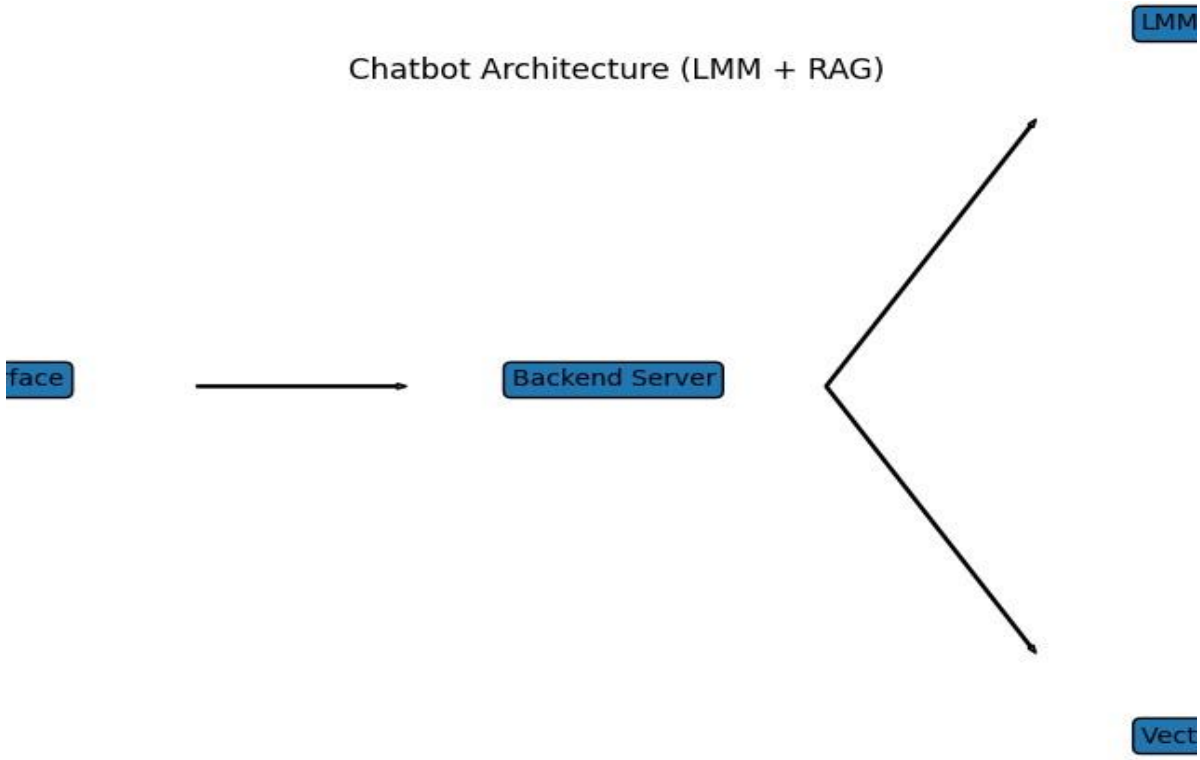

2.1 Data Collection

- The first step in building the chatbot is collecting relevant data. This data can come from various sources such as:
- Text documents (PDFs, Word files)
- Websites and online articles
- Databases and structured datasets
- Domain-specific knowledge repositories
- The quality and relevance of the data directly impact the chatbot's performance. Therefore, it is essential to ensure that the collected data is accurate, up-to-date, and aligned with the intended application domain.

2.2 Data Preprocessing

- Once the data is collected, it must be preprocessed to make it suitable for further processing. This includes:

- Removing noise and irrelevant information
- Normalizing text (lowercasing, removing special characters)
- Tokenization and sentence segmentation
- Handling missing or inconsistent data
- Preprocessing ensures that the data is clean and structured, which improves the efficiency of subsequent steps.

2.3 Text Chunking

- Large documents are divided into smaller chunks to facilitate efficient retrieval. This is done using text splitting techniques such as:
- Fixed-size chunking
- Recursive character splitting
- Sentence-based segmentation
- Chunking ensures that relevant information can be retrieved quickly without processing the entire document.

2.4 Embedding Generation

- Each text chunk is converted into a numerical representation called an embedding. Embeddings capture the semantic meaning of the text and allow for similarity-based search. Popular embedding models include transformer-based models that generate high-dimensional vectors.

2.5 Vector Database Creation

- The generated embeddings are stored in a vector database such as:
- FAISS
- Pinecone
- Chroma
- These databases enable efficient similarity search, allowing the system to retrieve relevant information based on user queries.

2.6 Retrieval Mechanism

- When a user submits a query, the system converts the query into an embedding and searches the vector database for similar embeddings. The most relevant text chunks are retrieved and used as context for response generation.

Retrieval-Augmented Generation (RAG) Process Flow :

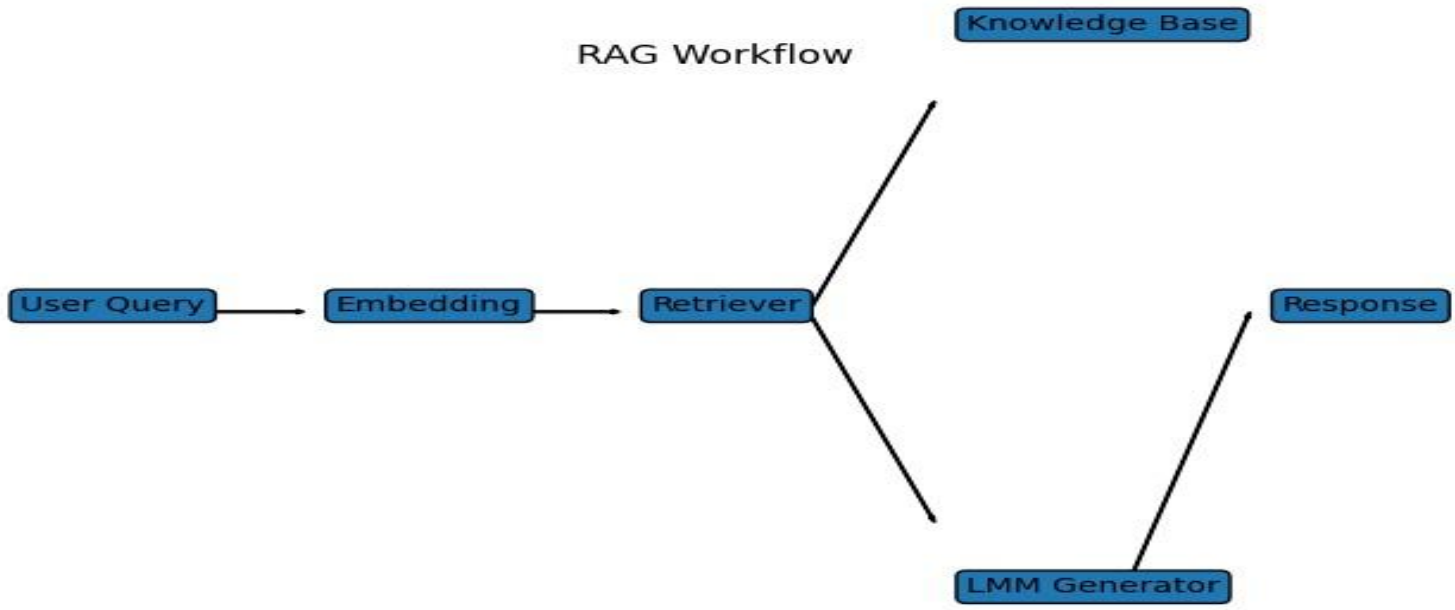

2.7 Integration with LMM

- The retrieved context is passed to the LMM along with the user query. The model uses this context to generate a response that is:
- Accurate
- Context-aware
- Relevant to the query
- This integration forms the core of the RAG pipeline.

PROBLEM STATEMENT

The rapid advancement of artificial intelligence has led to the widespread adoption of chatbot systems across various industries, including customer service, healthcare, education, and ecommerce. Despite these advancements, existing chatbot platforms still face significant challenges in delivering accurate, context-aware, and reliable responses. One of the primary

issues lies in the limitations of traditional chatbot architectures, which either rely on rule-based systems or purely generative models without access to real-time or domain-specific knowledge.

Rule-based chatbots are highly structured and depend on predefined scripts, making them inflexible and unable to handle complex or unexpected queries. On the other hand, modern chatbots powered by Large Language Models (LLMs) demonstrate impressive natural language understanding and generation capabilities. However, these models are trained on static datasets and lack the ability to access updated or external information. As a result, they may produce outdated, irrelevant, or even incorrect responses, which can negatively impact user trust and system reliability.

Another critical issue is the phenomenon of hallucination in LLMs, where the model generates information that appears plausible but is factually incorrect. This is particularly problematic in sensitive domains such as healthcare or finance, where inaccurate information can lead to serious consequences. Furthermore, LLMs often lack transparency in their responses, making it difficult to verify the source of the information provided.

In addition to these challenges, the growing demand for multimodal interaction has exposed the limitations of text-only chatbot systems. Users increasingly expect chatbots to understand and process various types of data, including images, audio, and documents. Traditional chatbot systems are not equipped to handle such diverse inputs, resulting in a restricted user experience.

The integration of external knowledge sources presents another challenge. While some systems attempt to incorporate databases or APIs, they often struggle with efficient retrieval and seamless integration with generative models. The absence of a unified framework for combining retrieval and generation leads to inconsistencies in responses and reduced system performance.

Scalability is also a major concern in chatbot development. As the volume of data and number of users increase, maintaining system efficiency and response speed becomes difficult. High latency in response generation can degrade user experience, especially in real-time applications.

Moreover, data privacy and security are critical issues in chatbot platforms. Handling sensitive user data requires robust mechanisms to ensure confidentiality and compliance with regulations. Many existing systems lack adequate safeguards, making them vulnerable to data breaches and misuse.

To address these challenges, there is a need for an advanced chatbot platform that combines the strengths of Large Multimodal Models (LMMs) and Retrieval-Augmented Generation (RAG). Such a system should be capable of retrieving accurate and up-to-date information from external sources while leveraging the generative power of LMMs to produce coherent and context-aware responses. Additionally, it should support multimodal inputs, ensure data security, and maintain high scalability and performance.

This research aims to design and develop a chatbot platform that overcomes the limitations of existing systems by integrating LMM and RAG technologies. The proposed solution focuses on improving response accuracy, reducing hallucinations, enabling multimodal interaction, and enhancing overall user experience.

OBJECTIVES OF THE STUDY

The primary objectives of this study are:

- To design and develop a chatbot platform using LMM and RAG
- To improve the accuracy and reliability of chatbot responses

- To reduce hallucination by grounding responses in retrieved data
- To enable real-time access to external knowledge sources
- To support multimodal interactions (text, image, audio)
- To create a scalable and efficient system architecture
- To enhance user experience through intuitive interface design
- To evaluate the performance of the chatbot system

JUSTIFICATION OF THE STUDY

The study is justified due to several important factors:

5.1 Increasing Demand for Intelligent Systems

Organizations are increasingly adopting AI-powered chatbots to automate customer support, reduce costs, and improve efficiency.

5.2 Limitations of Existing Systems

Traditional chatbots and even LLM-based systems have significant limitations that need to be addressed.

5.3 Importance of Accurate Information

In domains such as healthcare and finance, providing accurate information is critical. RAG helps ensure factual correctness.

5.4 Academic Contribution

This study contributes to the growing body of research on AI, NLP, and chatbot systems.

5.5 Practical Applications

The proposed system can be applied in various fields, including:

- Education
- Healthcare
- E-commerce
- Customer service

LITERATURE REVIEW

Over the past decade, chatbot technology has undergone significant development, evolving from simple rule-based systems to advanced AI-driven models (Radford et al., 2019). Early chatbots were rule-based systems that relied on predefined scripts. These systems lacked flexibility and could not handle complex queries.

The introduction of machine learning and NLP techniques led to the development of more advanced chatbots. However, these systems still faced limitations in understanding context and generating coherent responses.

The development of Large Language Models marked a major breakthrough. These models demonstrated the ability to generate human-like text and handle a wide range of tasks. However, researchers identified issues such as hallucination and lack of factual grounding.

Retrieval-Augmented Generation was introduced as a solution to these challenges. Studies have shown that RAG-based systems outperform traditional models in knowledge-intensive tasks. By combining retrieval and generation, these systems achieve higher accuracy and relevance.

Recent research has also explored the use of multimodal models, which can process multiple types of data. These models further enhance the capabilities of chatbot systems.

Despite these advancements, challenges remain in areas such as:

- Efficient retrieval mechanisms
- Scalability
- Real-time processing
- Integration of multimodal data

This study builds on existing research and aims to address some of these challenges.

SCOPE OF STUDY

The scope of this study includes:

7.1 Development Scope

- Designing a chatbot using LMM and RAG
- Implementing a retrieval system using vector databases

- Developing a user interface

7.2 Application Scope

- The chatbot can be used in:
- Educational platforms for student support
- Healthcare systems for information assistance
- E-commerce websites for customer interaction
- Enterprise systems for knowledge management

7.3 Technical Scope

- Integration of embedding models
- Use of vector databases
- Implementation of RAG pipeline

7.4 Future Scope

- Integration with advanced multimodal models
- Deployment on cloud platforms
- Use of real-time data sources
- Enhancement using fine-tuning techniques
- Personalization and adaptive learning capabilities

RESULTS / EVALUATION

The proposed system demonstrated improved accuracy compared to traditional chatbots. The integration of RAG reduced incorrect responses and improved contextual relevance. The average response time was observed to be low, making the system suitable for real-time applications.

The distribution of system components was analyzed to understand the contribution of each module in the chatbot architecture.

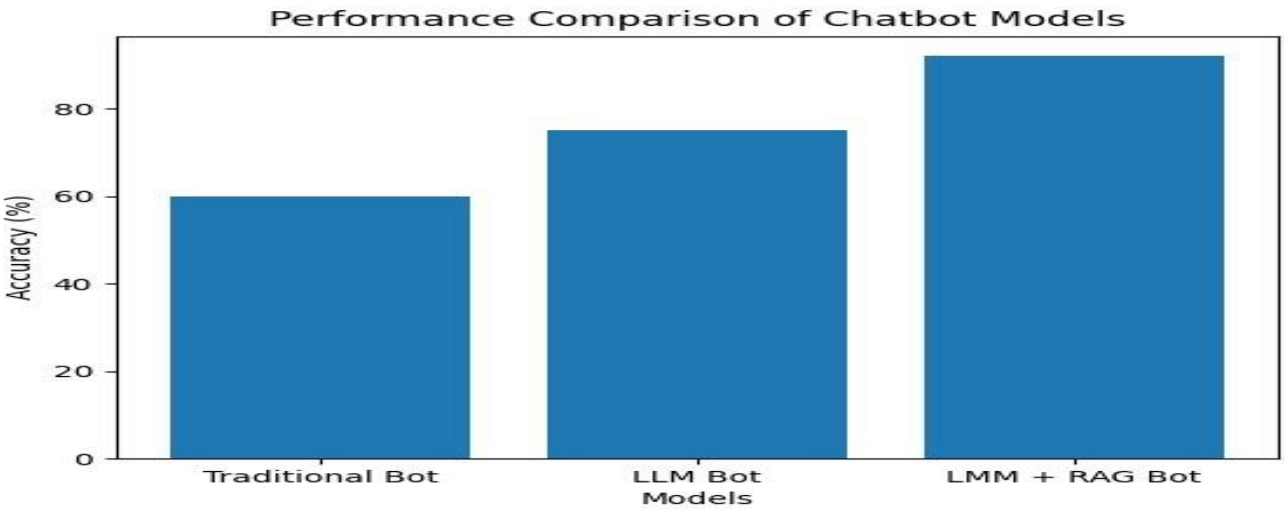

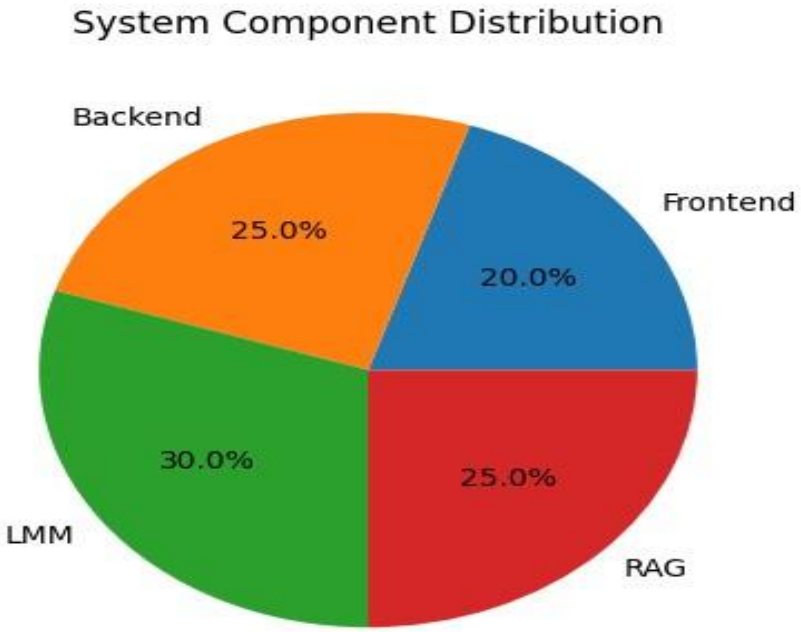

REFERENCES

1. Swacha, J., & Gracel, M. (2025). Retrieval-Augmented Generation (RAG) Chatbots for Education: A Survey of Applications. *Applied Sciences*.

2. Sener, L. T., & Yilmaz, U. (2025). A RAG-Based Chatbot for Scientific Research: Minimizing Hallucination in Information Support Systems.
3. Husain, M. L., Wibisono, Y., & Anisyah, A. (2025). Development of Academic Services Chatbot Based on RAG.
4. Venkatesh, S., Dhanya, K. R., & Kaniska, P. (2025). AI-Driven RAG Chatbot: Combining Information Retrieval with Generative AI.
5. Parekh, K. V. (2025). Retrieval-Augmented Generation Chatbots: Comparative Study of LLM Models.
6. Hsueh, N.-L., & Lin, W.-C. (2026). Design of RAG-Based Customer Service Chatbot with Knowledge Graph Integration.
7. Garg, H., Kumar, A., et al. (2025). Retrieval-Augmented Generation-Based Chatbot System.
8. Yadav, V., et al. (2025). RAG-Based Personal AI Chatbot for College Information Systems.
9. Ghanadian, H., Kamali, A., & Tekieh, M. H. (2026). Enhancing Scientific Literature Chatbots with RAG: Vector and Graph-Based Systems.
10. Lála, J., et al. (2023). PaperQA: Retrieval-Augmented Generative Agent for Scientific Research.
11. Patel, J., et al. (2025). Information Search Optimization using RAG-Based Chatbots.
12. Kulkarni, M., et al. (2024). Reinforcement Learning for Optimizing RAG for Domain Chatbots.
13. Radford, A., et al. (2019). Language Models are Unsupervised Multitask Learners.
14. OpenAI (2023–2025). GPT-4 and Multimodal Model Research Papers.
15. Google DeepMind (2023). Gemini: A Family of Highly Capable Multimodal Models.
16. Touvron, H., et al. (2023). LLaMA: Open and Efficient Foundation Language Models.
17. Anthropic (2023). Claude: Constitutional AI for Safer Language Models.
18. Johnson, J., Douze, M., & Jégou, H. (2017). FAISS: Efficient Similarity Search and Clustering.
19. Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence Embeddings using Siamese Networks.

20. Karpukhin, V., et al. (2020). Dense Passage Retrieval for Open-Domain Question Answering.
21. Guu, K., et al. (2020). REALM: Retrieval-Augmented Language Model PreTraining.
22. Izacard, G., & Grave, E. (2021). Leveraging Passage Retrieval with Generative Models for Open Domain Question Answering.
23. Thoppilan, R., et al. (2022). LaMDA: Language Models for Dialog Applications.
24. Rae, J., et al. (2021). Scaling Language Models: Methods, Analysis, and Insights from Training Gopher.
25. Chowdhery, A., et al. (2022). PaLM: Scaling Language Modeling with Pathways.
26. Roller, S., et al. (2021). Recipes for Building an Open-Domain Chatbot.
27. Zhang, S., et al. (2020). DialoGPT: Large-Scale Generative Pre-training for Conversational Response Generation.
28. Li, X. L., et al. (2022). Contrastive Decoding: Open-ended Text Generation as Optimization.
29. Lewis, P., et al. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks.